Imfit documentation
Release 0.1

Matthew Newville

August 12, 2011

CONTENTS

The lmfit Python package provides a simple, flexible interface to non-linear least-squares optimization, or curve fitting. By default, LMFIT uses the Levenberg-Marquardt minimization algorithm from MINPACK-1 as implemented in scipy.optimize.leastsq, but it can also use the L-BFGS (limited memory Broyden-Fletcher-Goldfarb-Shanno) algorithm as implemented in scipy.optimize.l_bfgs_b or the simulated annealing algorithm as implemented in scipy.optimize.anneal. Support for other optimization routines may be added in the future. While these functions from scipy.optimize provide the core numerical algorithm for non-linear least-squares minimization, the lmfit package adds a few simple conveniences. Most of this document will assume that the Levenberg-Marquardt algorithm is being discussed, as it appears to be the most robust for finding local minima of well-described models of scientific measurements.

For any minimization problem, the programmer must provide a function that takes a set of values for the variables in the fit, and produces the residual function to be minimized in the least-squares sense.

The lmfit package allows models to be written in terms of a set of Parameters, which are extensions of simple numerical variables with the following properties:

- Parameters can be fixed or floated in the fit.

- Parameters can be bounded with a minimum and/or maximum value.

- Parameters can be written as simple mathematical expressions of other Parameters. These values will be re-evaluated at each step in the fit, so that the expression is statisfied. This gives a simple but flexible approach to constraining fit variables.

The main advantage to using Parameters instead of fit variables is that the model function does not have to be rewritten for a change in what is varied or what constraints are placed on the fit. The programmer can write a fairly general model, and allow a user of the model to change what is varied and what constraints are placed on the model.

In addition, lmfit calculates are reports the estimated uncertainties and correlation between fitted variables.

DOWNLOADING AND INSTALLATION

1.1 Prerequisites

The lmfit package requires Python, Numpy, and Scipy. Extensive testing on version compatibility has not yet been done. I have not yet tested with Python 3.

1.2 Downloads

The latest stable version is available from a few sources:

Download Option	Location
Source Kit	lmfit-0.1.tar.gz (CARS) or
Windows Installers	lmfit-0.1.win32-py2.6.exe or lmfit-0.1.win32-py2.7.exe
Development Version	use lmfit github repository

if you have Python Setup Tools installed, you can download and install the PyEpics Package simply with:

```
easy_install -U lmfit
```

1.3 Development Version

To get the latest development version, use:

```
git clone http://github.com/newville/lmfit-py.git
```

1.4 Installation

Installation from source on any platform is:

```
python setup.py install
```

1.5 License

The LMFIT-py code is distribution under the following license:

GETTING STARTED WITH NON-LINEAR LEAST-SQUARES FITTING

The lmfit package is designed to provide a simple way to build complex fitting models and apply them to real data. This chapter describes how to set up and perform simple fits, but does assume some basic knowledge of Python, Numpy, and modeling data.

To model data in the least-squares sense, the most important step is writing a function that takes the values of the fitting variables and calculates a residual function (data-model) that is to be minimized in the least-squares sense

$$\chi^2 = \sum_i^N \frac{[y_i^{\mathrm{meas}} - y_i^{\mathrm{model}}(\mathbf{v})]^2}{\epsilon_i^2}$$

where y_i^{meas} is the set of measured data, $y_i^{\mathrm{model}}(\mathbf{v})$ is the model calculation, $\mathbf{v}$ is the set of variables in the model to be optimized in the fit, and ϵ_i is the estimated uncertainty in the data.

In a traditional non-linear fit, one writes a function that takes the variable values and calculates the residual $y_i^{\mathrm{meas}} - y_i^{\mathrm{model}}(\mathbf{v})$, perhaps something like:

```python
def residual(vars, x, data):
    amp = vars[0]
    phaseshift = vars[1]
    freq = vars[2]
    decay = vars[3]

    model = amp * sin(x * freq  + phaseshift) * exp(-x*x*decay)

    return (data-model)
```

To perform the minimization with scipy, one would do:

```python
from scipy.optimize import leastsq
vars = [10.0, 0.2, 3.0, 0.007]
out = leastsq(residual, vars, args=(x, data))
```

Though in python, and fairly easy to use, this is not terribly different from how one would do the same fit in C or Fortran.

2.1 Using `Parameters` instead of Variables

As described above, there are several practical challenges in doing least-squares fit with the traditional implementation (Fortran, scipy.optimize.leastsq, and most other) in which a list of fitting variables to the function to be minimized. These challenges include:

1. The user has to keep track of the order of the variables, and their meaning – vars[2] is the frequency, and so on.

2. If the user wants to fix a particular variable (*not* vary it in the fit), the residual function has to be altered. While reasonable for simple cases, this quickly becomes significant work for more complex models, and greatly complicates modeling for people not intimately familiar with the code.

3. There is no way to put bounds on values for the variables, or enforce mathematical relationships between the variables.

LMFIT is designed to void these shortcomings.

The main idea of LMFIT is to expand a numerical variable with a `Parameter`, which have more attributes than simply their value. Instead of a pass a list of numbers to the function to minimize, you create a `Parameters` object, add parameters to this object, and pass along this object to your function to be minimized. With this transformation, the above example would be translated to look like:

```python
from lmfit import minimize, Parameters

def residual(params, x, data):
    amp = params['amp'].value
    pshift = params['phase'].value
    freq = params['frequency'].value
    decay = params['decay'].value

    model = amp * sin(x * freq  + pshift) * exp(-x*x*decay)

    return (data-model)

params = Parameters()
params.add('amp', value=10)
params.add('decay', value=0.007)
params.add('phase', value=0.2)
params.add('frequency', value=3.0)

out = minimize(residual, params, args=(x, data))
```

So far, this simply looks like it replaced a list of values with a dictionary, accessed by name. But each of the named `Parameter` in the `Parameters` object hold additional attributes to modify the value during the fit. For example, Parameters can be fixed or bounded, and this can be done when being defined:

```python
params = Parameters()
params.add('amp', value=10, vary=False)
params.add('decay', value=0.007, min=0.0)
params.add('phase', value=0.2)
params.add('frequency', value=3.0, max=10)
```

or later:

```python
params['amp'].vary = True
params['decay'].max = 0.10
```

Now the fit will *not* vary the amplitude parameter, and will also impose a lower bound on the decay factor and an upper bound on the frequency. Importantly, our function to be minimized remains unchanged.

An important point here is that the *params* object can be copied and modified to make many user-level changes to the model and fitting process. Of course, most of the inormation about how your data is modeled goes into the fitting function, but the approach here allows some external control as well.

2.2 The `Parameter` class

class `Parameter` (*value=None*[, *vary=True*[, *min=None*[, *max=None*[, *name=None*[, *expr=None*]]]]])
    create a Parameter object. These are the fundamental extension of a fit variable within LMFIT, but you will probably create most of these with the `Parameters` class.

      Parameters

- **value** – the numerical value for the parameter

- **vary** (boolean (`True`/`False`)) – whether to vary the parameter or not.

- **min** – lower bound for value (`None` = no lower bound).

- **max** – upper bound for value (`None` = no upper bound).

- **name** (`None` or string – will be overwritten during fit if `None`.) – parameter name

- **expr** (`None` or string) – mathematical expression to use to evaluate value during fit.

Each of these inputs is turned into an attribute of the same name. As above, one hands a dictionary of Parameters to the fitting routines. The name for the Parameter will be set to be consistent

After a fit, a Parameter for a fitted variable (ie with vary = `True`) will have the `value` attribute holding the best-fit value, and may (depending on the success of the fit) have obtain additional attributes.

`stderr`
    the estimated standard error for the best-fit value.

`correl`
    a dictionary of the correlation with the other fitted variables in the fit, of the form:

```
{'decay': 0.404, 'phase': -0.020, 'frequency': 0.102}
```

The `expr` attribute can contain a mathematical expression that will be used to compute the value for the Parameter at each step in the fit. See *Using Mathematical Constraints* for more details and examples of this feature.

2.3 The `Parameters` class

class `Parameters`
    create a Parameters object. This is little more than a fancy dictionary, with the restrictions that

1. keys must be valid Python symbol names (so that they can be used in expressions of mathematical constraints). This means the names must match `[a-z_][a-z0-9_]*` and cannot be a Python reserved word.

    2.values must be vaild `Parameter` objects.

Two methods for provided for convenience of initializing Parameters.

add (*name*[, *value=None*[, *vary=True*[, *min=None*[, *max=None*[, *expr=None*]]]]])
    add a named parameter. This simply creates a `Parameter` object associated with the key *name*, with optional arguments passed to `Parameter`:

```
p = Parameters()
p.add('myvar', value=1, vary=True)
```

add_many (*self*, *paramlist*)
    add a list of named parameters. Each entry must be a tuple with the following entries:

```
name, value, vary, min, max, expr
```

That is, this method is somewhat rigid and verbose (no default values), but can be useful when initially defining a parameter list so that it looks table-like:

```
p = Parameters()
#            (Name,  Value,  Vary,   Min,  Max,   Expr)
p.add_many(('amp1',    10,   True, None, None,   None),
           ('cen1',    1.2,  True,  0.5,  2.0,   None),
           ('wid1',    0.8,  True,  0.1, None,   None),
           ('amp2',    7.5,  True, None, None,   None),
           ('cen2',    1.9,  True,  1.0,  3.0,   None),
           ('wid2',  None, False, None, None, '2*wid1/3'))
```

2.4 Simple Example:

Putting it all together, a simple example of using a dictionary of `Parameter`s and `minimize()` might look like this:

```python
from lmfit import minimize, Parameters

def residual(params, x, data=None):
    amp = params['amp'].value
    shift = params['phase_shift'].value
    omega = params['omega'].value
    decay = params['decay'].value

    model = amp * sin(x * omega + shift) * exp(-x*x*decay)

    return (data-model)

params = Parameters()
params.add('amp', value=10)
params.add('decay', value=0.007, vary=False)
params.add('phase_shift', value=0.2)
params.add('omega', value=3.0)

result = minimize(residual, params, args=(x, data))

print result.chisqr
print 'Best-Fit Values:'
for name, par in params.items():
    print '  %s = %.4f +/- %.4f ' % (name, par.value, par.stderr)
```

The results of this fit looks like this:

```
OUTPUT HERE...
```

PERFORMING FITS, ANALYZING OUTPUTS

As shown in the previous sections, a simple fit can be performed with the `minimize()` function. For more sophisticated modeling, the `Minimizer` class can be used to gain a bit more control, especially when using complicated constraints.

3.1 The `minimize()` function

The minimize function takes a function to minimze, a dictionary of `Parameter` , and several optional arguments. See *Writing a Fitting Function* for details on writing the function to minimize.

minimize (*function*, *params*[, *args=None*[, *kws=None*[, *engine='leastsq'*[, ***leastsq_kws*]]]])
find values for the params so that the sum-of-squares of the returned array from function is minimized.

> **Parameters**
>
> - **function** (*callable.*) – function to return fit residual. See *Writing a Fitting Function* for details.
> - **params** (*dict*) – a dictionary of Parameters. Keywords must be strings that match [a-z_][a-z0-9_]* and is not a python reserved word. Each value must be `Parameter`.
> - **args** (*tuple*) – arguments tuple to pass to the residual function as positional arguments.
> - **kws** (*dict*) – dictionary to pass to the residual function as keyword arguments.
> - **engine** (*string*) – name of fitting engine to use. See *Choosing Different Fitting Engines* for details
> - **leastsq_kws** (*dict*) – dictionary to pass to scipy.optimize.leastsq
>
> **Returns** Minimizer object, which can be used to inspect goodness-of-fit statistics, or to re-run fit.

On output, the params will be updated with best-fit values and, where appropriate, estimated uncertainties and correlations. See *Goodness-of-Fit and estimated uncertainty and correlations* for further details.

3.2 Writing a Fitting Function

An important component of a fit is writing a function to be minimized in the least-squares sense. Since this function will be called by other routines, there are fairly stringent requirements for its call signature and return value. In

principle, your function can be any python callable, but it must look like this:

func(params, *args, **kws):
> calculate residual from parameters.

>> **Parameters**

>>> • **params** (*dict*) – parameters.

>>> • **args** – positional arguments. Must match *args* argument to `minimize()`

>>> • **kws** – keyword arguments. Must match *kws* argument to `minimize()`

>> **Returns** residual array (generally data-model) to be minimized in the least-squares sense.

>> **Return type** numpy array. The length of this array cannot change between calls.

A common use for the positional and keyword arguments would be to pass in other data needed to calculate the residual, including such things as the data array, dependent variable, uncertainties in the data, and other data structures for the model calculation.

As the function will be passed in a dictionary of `Parameter`s, it is advisable to unpack these to get numerical values at the top of the function. A simple example would look like:

```python
def residual(pars, x, data=None):
    # unpack parameters:
    #  extract .value attribute for each parametr
    amp = pars['amp'].value
    period = pars['period'].value
    shift = pars['shift'].value
    decay = pars['decay'].value

    if abs(shift) > pi/2:
        shift = shift - sign(shift)*pi

    if abs(period) < 1.e-10:
        period = sign(period)*1.e-10

    model = amp * sin(shift + x/per) * exp(-x*x*decay*decay)

    if data is None:
        return model
    return (model - data)
```

In this example, x is a positional (required) argument, while the `data` array is actually optional (so that the function returns the model calculation if the data is neglected). Also note that the model calculation will divide x by the varied value of the 'period' Parameter. It might be wise to make sure this parameter cannot be 0. It would be possible to use the bounds on the `Parameter` to do this:

```python
params['period'] = Parameter(value=2, min=1.e-10)
```

but might be wiser to put this directly in the function with:

```python
if abs(period) < 1.e-10:
    period = sign(period)*1.e-10
```

3.3 Choosing Different Fitting Engines

By default, the Levenberg-Marquardt algorithm is for fitting. While often criticized as finding only *local* minima, this approach has some distinct advantages, in that it is fast, well-behaved for most curve-fitting needs, and makes it

easy to estimate uncertainties for and correlations between pairs of fit variables, as discussed in *Goodness-of-Fit and estimated uncertainty and correlations*.

Alternative algorithms can also be used.

3.4 Goodness-of-Fit and estimated uncertainty and correlations

On a successful fit using the *leastsq* engine, several goodness-of-fit statistics and values related to the uncertainty in the fitted variables will be calculated. These are all encapsulated in the `Minimizer` object for the fit, as returned by `minimize()`. The values related to the entire fit are stored in attributes of the `Minimizer` object, as shown in *Table of Goodness-of-Fit Statistics* while those related to each fitted variables are stored as attributes of the corresponding `Parameter`.

Table of Goodness-of-Fit Statistics: These statistics are all attributes of the `Minimizer` object returned by `minimize()`.

Minimizer Attribute	Description / Formula
nfev	number of function evaluations
success	boolean (`True`/`False`) for whether fit succeeded.
errorbars	boolean (`True`/`False`) for whether uncertainities were estimated.
message	message about fit success.
ier	integer error value from scipy.optimize.leastsq
lmdif_message	message from scipy.optimize.leastsq
nvarys	number of variables in fit N_{varys}
ndata	number of data points: N
nfree	degrees of freedom in fit: $N - N_{\text{varys}}$
residual	residual array (return of `func()`: Resid
chisqr	chi-square: $\chi^2 = \sum_i^N [\text{Resid}_i]^2$
redchi	reduced chi-square: $\chi_\nu^2 = \chi^2/(N - N_{\text{varys}})$

Note that the calculation of chi-square and reduced chi-square assume that the returned residual function is scaled properly to the uncertainties in the data. For these statistics to be meaningful, the person writing the function to function to be minimized must scale them properly.

When possible, standard errors for the fitted variables, and correlations between pairs of fitted variables are automatically calculated after the fit is performed. The standard error (estimated 1σ error-bar) go into the `stderr` attribute of the Parameter. The correlations with all other variables will be put into the `correl` attribute of the Parameter – a dictionary with keys for all other Parameters and values of the corresponding correlation.

In some cases, it may not be possible to estimate the errors and correlations. For example, if a variable actually has no practical effect on the fit, it will likely cause the covariance matrix to be singular, so that errors cannot be estimated. Placing bounds on varied Parameters makes it more likely that errors cannot be estimated, as being near the maximum or minimum value makes the covariance matrix singular. In these cases, the `errorbars` attribute of the fit result (`Minimizer` object) will be `False`.

3.5 Using the `Minimizer` class

For full control of the fitting process, you'll want to create a `Minimizer` object, or at least use the one returned from the `minimize()` function.

class **Minimizer** (*function params*$\big[$, *fcn_args=None*$\big[$, *fcn_kwsn=None*$\big[$, *engine='leastsq'*$\big[$, ***kws*$\big]\big]\big]\big]$)
  creates a Minimizer object.

    Parameters

- **function** (*callable.*) – function to return fit residual. See *Writing a Fitting Function* for details.

- **params** (*dict*) – a dictionary of Parameters. Keywords must be strings that match `[a-z_][a-z0-9_]*` and is not a python reserved word. Each value must be `Parameter`.

- **fcn_args** (*tuple*) – arguments tuple to pass to the residual function as positional arguments.

- **fcn_kws** (*dict*) – dictionary to pass to the residual function as keyword arguments.

- **leastsq_kws** (*dict*) – dictionary to pass to scipy.optimize.leastsq

Returns Minimizer object, which can be used to inspect goodness-of-fit statistics, or to re-run fit.

The Minimizer object has only a few methods.

USING MATHEMATICAL CONSTRAINTS

While being able to fix variables and place upper and lower bounds on their values are key parts of LMFIT, an equally important feature is the ability to place mathematical constraints on parameters. This section describes how to do this, and what sort of parameterizations are possible.

4.1 Overview

Just as one

4.2 Supported Operators, Functions, and Constants

The mathematicl expressions used to define constrained Parameters need to be valid python expressions. The full set of built-in

4.3 Some Implementation Details

4.4 Advanced usage – beyond LMFIT

INDEX